

CHATGPT UNLEASHED



TODAY'S JOURNEY

Today we move from individual conversations to structured, repeatable AI workflows. We cover connected applications, Projects in depth, reusable Skills, and multi-stage workflows that combine research, synthesis, generation, and review.



DAY 2 MODULES AT A GLANCE

Today covers six content modules and four labs:

- **Module A** — From Conversations to Workflows
- **Module B** — Connected Applications and Tools
- **Module C** — Projects in Depth
- **Module D** — Skills and Reusable AI Workflows
- **Module E** — Advanced ChatGPT Workflows
- **Labs 5–8** — Connected Apps, Advanced Workflows, Reusable Workflow, Reusable Skill

RECAP: WHAT WE COVERED YESTERDAY

Day 1 established the foundation for everything that follows today:

- How LLMs generate responses: tokens, context windows, and nondeterministic output
- Structured prompting: goal, context, constraints, format, and audience
- Document upload, image analysis, CSV data analysis, and web search
- ChatGPT Projects as persistent workspaces and deep research as a dedicated mode

HOW TODAY BUILDS ON DAY 1

Yesterday you learned how to communicate effectively with ChatGPT as an individual. Today you learn how to organize that communication into repeatable work your entire team can use.

A great prompt is a one-time tool. A workflow is an organizational asset.



THE PROBLEM WITH ONE-OFF CHATS

One-off chats are fast but disposable. You get a useful response, close the tab, and the next time you need the same type of output you start over from scratch.

The hidden cost compounds: repeated setup, inconsistent output quality, and no institutional memory of what worked last time.

WHAT MAKES A PROCESS A WORKFLOW?

A workflow is a repeatable sequence of steps that takes defined inputs and produces a defined output. Each step has a clear purpose, a specific intermediate artifact, and a clear handoff to the next step.

Converting a successful chat sequence into a workflow means capturing the prompts, the checkpoints, and the decisions — not just the final result.



DESIGNING BEFORE YOU PROMPT

The most common workflow mistake is opening ChatGPT before answering three questions: what am I making, what inputs do I need, and where must a human review the output before we continue?

Design on paper first. The prompt comes second.

DECOMPOSING A TASK INTO STAGES

Complex outputs should not come from a single prompt. Breaking a task into stages produces better results because each stage can be reviewed before the next begins.

Typical stages: research, synthesis, drafting, critique, refinement. Each stage has its own prompt and its own intermediate output.



WHAT EACH STAGE SHOULD PRODUCE

Every stage in a workflow should produce a specific, reviewable artifact before the next stage begins.


- Research → a bullet-point findings list with sources
- Synthesis → a structured outline, not prose
- Draft → the full document written against the approved outline
- Critique → a numbered list of specific issues
- Refinement → the revised document with changes noted

The artifact is what you review — not a vague sense that the conversation is going well.

HUMAN-IN-THE-LOOP WORKFLOWS

Not every step should run to completion without human review. The right design includes explicit checkpoints where a person evaluates the output and decides whether to continue, revise, or escalate.

Skipping checkpoints does not save time — it transfers the risk of bad output to whoever reads the final deliverable.



APPROVAL POINTS: WHAT TO REVIEW


At each checkpoint, answer three questions: does this output match the specification, are there claims I cannot verify, and is this ready to become input to the next stage?

If the answer to any question is no, correct the output before continuing. Errors in intermediate outputs compound through every stage that follows.

WHAT TO AUTOMATE AND WHAT TO KEEP HUMAN

Automation suits tasks that are repetitive, have fixed output formats, and carry low consequences for error. Human review is essential when output will be presented to leadership, acted on financially, or attributed to your organization.

The default is human-in-the-loop until you have evidence that a step is reliable enough to run unsupervised.



CONTEXT CONTINUITY ACROSS STEPS


Context from earlier steps informs later steps. In a multi-stage workflow, ensure ChatGPT has the relevant output from the previous stage before asking it to perform the next task.

Do not assume ChatGPT remembers prior output from a previous session. Stay in the same conversation or paste the prior output explicitly when starting a new one.

BREAKING COMPLEX PROBLEMS INTO SMALLER TASKS

The same principle that applies to workflow stages applies to individual prompts: one task per message produces better output than combining multiple tasks.

"Summarize this document, extract the key risks, and draft a recommendation" is three tasks. Run them separately and review each output before continuing.



HAND-OFF A PLAN

From Complex Request to Results

Plan it. Pass it. Execute it.

Break it down
to get it done.



Specialized Agents

Use the right agent for the right job — one plans, another executes.



Better Results

Structured plans lead to higher quality, more consistent outcomes.



More Control

You can review, adjust, or rerun the plan at any time.

THE COMPOUND EFFECT OF GOOD INTERMEDIATE OUTPUTS

When each stage produces a clean, reviewed output, each subsequent stage starts from a stronger foundation. Errors caught at the research stage cost one revision. Errors caught at the draft stage cost five.

The return on reviewing intermediate outputs is always positive.

MANAGING STATE BETWEEN WORKFLOW SESSIONS

Not all workflows complete in a single session. When a workflow spans multiple conversations, the analyst is responsible for carrying forward the reviewed output from each prior stage.

Copy the approved intermediate artifact into the next conversation explicitly. ChatGPT does not retrieve prior sessions automatically.



WORKFLOW FAILURE MODES: WHAT GOES WRONG

Three patterns that reliably break multi-stage workflows:

- Asking for the final deliverable before collecting raw material
- Moving to the next stage without reviewing the current one
- Writing templates so vague that different analysts get different results from the same inputs

From Ideas to a Better Result

A human-in-the-loop workflow for higher quality output




WORKFLOW DESIGN CHECKLIST

Before building any workflow, verify you can answer these questions:

- What exact inputs must the analyst provide before starting?
- What does each stage produce before the next begins?
- Which checkpoints require human approval, and what is being reviewed?
- What does the final deliverable look like: format, length, and audience?

A workflow you cannot describe without running it is not a workflow — it is a conversation.



WHAT ARE CONNECTED APPLICATIONS?


ChatGPT can connect to external services: cloud storage (Google Drive, Dropbox, OneDrive), productivity tools, and other approved integrations. These connections allow ChatGPT to retrieve content from those sources and, in some cases, take actions such as creating files or sending messages.

Connections fundamentally change what ChatGPT can access — and what can go wrong.

RETRIEVAL VS. ACTION: A CRITICAL DISTINCTION

There are two types of connected capabilities. Retrieval means reading — finding a file and summarizing it. Action means changing something — creating a file, sending a message, updating a record.

A retrieval error produces a bad summary. An action error sends an email or overwrites a record. These carry very different consequences.



ACTIONS VS. RETRIEVAL: THE RISK MULTIPLIER

For every action capability, define the failure you want to prevent and the control that prevents it before enabling it.


- Create a file → requires folder scoping; no overwrite without confirmation
- Send a message → requires explicit user approval before the message is sent
- Update a record → requires human review of the proposed change, not just the intent

"The user can undo it" is not a sufficient control for any of these.

WHAT HAPPENS WHEN CHATGPT CONNECTS TO YOUR DATA

When ChatGPT retrieves content from a connected application, that content enters the conversation. It moves from the original system's access controls into OpenAI's infrastructure.

This is not necessarily wrong — but the data is no longer protected by the source system's controls. That is the boundary worth understanding before you authorize any connection.



DATA BOUNDARY: CONTENT LEAVES THE SOURCE SYSTEM

Connecting ChatGPT to Google Drive is not the same as using Google Drive. The moment ChatGPT retrieves a file, its content becomes part of the AI session.

The question is not just "can ChatGPT access this?" but "should this content leave the system where it is currently protected?"

PERMISSIONS AND LEAST PRIVILEGE

The principle of least privilege applies to AI connections. Grant the minimum access needed for the specific task. Scope it to the right folders, the right data types, and the right duration.

"Read access to all folders" is almost never the right permission for a specific use case.



FILE UPLOAD VS. CONNECTED APP: WHEN TO USE WHICH

File upload is always scoped — you control exactly which document enters the session, and nothing else. A connected app may access more than you intended, depending on how permissions were configured.

Use file upload when precision matters. Use connected apps when retrieval automation justifies the governance overhead it requires.

REVOKING ACCESS: THE PLAN YOU NEED BEFORE YOU CONNECT

Before authorizing a connection, know how to revoke it. Who can revoke it? How quickly? What data remains in the session after revocation?

If you cannot answer these questions before connecting, you have accepted a risk you have not bounded.



EVALUATING CONNECTION RISK

Before enabling a connection, classify the data that could be accessed. Public reports and drafts you own are generally fine. PII, legal documents, HR records, and financial data are not.

Classification happens before the connection is authorized — not after something goes wrong.

AUTHORIZATION FRAMEWORK FOR AI CONNECTIONS

Three questions to answer before authorizing any connected application:

- What is the minimum permission this specific task actually requires?
- What data in the source system should never appear in a ChatGPT session?
- What is the revocation process if something goes wrong?

If you cannot answer all three, do not authorize the connection yet.

LAB 5



POP QUIZ: Which scenario is generally acceptable without special approval?

- A. Connecting ChatGPT to the company's entire shared Google Drive for team-wide access.
- B. Uploading a public report you downloaded to your desktop and summarizing it.
- C. Allowing ChatGPT to update CRM records based on conversation output.
- D. Retrieving HR salary data to analyze compensation trends.



POP QUIZ: Which scenario is generally acceptable without special approval?

- A. Connecting ChatGPT to the company's entire shared Google Drive for team-wide access.
- B. Uploading a public report you downloaded to your desktop and summarizing it.**
- C. Allowing ChatGPT to update CRM records based on conversation output.
- D. Retrieving HR salary data to analyze compensation trends.



POP QUIZ: An analyst connects ChatGPT to Google Drive. What is the primary data risk?

- A. The content retrieved leaves the source system's access controls and enters the AI session.
- B. ChatGPT might misformat the document summary it produces.
- C. Google Drive will charge extra for third-party AI integrations.
- D. ChatGPT's context window cannot hold large documents retrieved from Drive.



POP QUIZ: An analyst connects ChatGPT to Google Drive. What is the primary data risk?


- A. The content retrieved leaves the source system's access controls and enters the AI session.**
- B. ChatGPT might misformat the document summary it produces.
- C. Google Drive will charge extra for third-party AI integrations.
- D. ChatGPT's context window cannot hold large documents retrieved from Drive.



PROJECTS AS PERSISTENT WORKSPACES

A ChatGPT Project is more than a saved conversation. It is a persistent workspace where instructions, uploaded files, and context accumulate and apply automatically to every session opened inside it.

The persistent nature of Projects makes them fundamentally different from one-off chats — and fundamentally more valuable for recurring work.



PROJECT INSTRUCTIONS VS. CONVERSATION CONTEXT


Project instructions apply to every conversation in the project automatically. Conversation-level context applies only within that specific session and disappears when a new conversation begins.

For teams, the project instruction is the one authoritative place where shared persona, output standards, and background context live.

COMBINING FILES, INSTRUCTIONS, AND CONVERSATIONS

A well-configured Project has three layers working together: project instructions that define context and style, uploaded files that provide factual grounding, and conversations that apply both to specific tasks.

Each layer has a distinct role. None replaces the other — they compose into something stronger than any one layer alone.



WHEN TO START A NEW PROJECT VS. A NEW CONVERSATION


Start a new conversation when the task is a variation on work already in the project. Start a new project when the topic, client, or context is genuinely separate from everything else.

Mixing unrelated work in one project creates context bleed — instructions and files intended for one task inadvertently shaping another.

BUILDING PROJECT TEMPLATES FOR TEAMS

A project template is a pre-configured workspace with standard instructions and common files already loaded. When a recurring task begins or a new analyst joins the team, they open the template rather than configuring from scratch.

The value is consistency: every analyst starts from the same baseline every time.



MANAGING MULTIPLE PROJECTS

Most professionals accumulate multiple Projects over time: one per client, one per recurring initiative, one for personal use. Keeping them separate prevents context bleed — the wrong persona or files appearing in the wrong conversation.

Separate projects are better than a combined one. Context bleed is subtle and hard to catch after it has already shaped the output.

NAMING AND ORGANIZING PROJECTS FOR LONG-TERM USE

Project names should describe the work unambiguously: "Apex Retail — AI Customer Service Q3 2025" rather than "AI Project." When you have twenty projects, clear names prevent opening the wrong one.

Naming standards are a lightweight form of governance — they make the workspace navigable for you and for teammates who inherit your work.



CONTEXT SEPARATION: WHY IT MATTERS

Context from one project should not appear in another. Client A's data should not surface when working on Client B's analysis. Personal notes should not inform a professional deliverable.

Projects enforce this boundary structurally. You cannot accidentally leak context between separate projects — the separation is built into the architecture.

WHAT IS A SKILL?

A Skill is a portable, saved set of instructions, templates, and supporting resources. Unlike project instructions scoped to a single project, Skills can be invoked in any conversation.

Skills are best for recurring tasks with a defined structure and quality standards that must be applied consistently every time the task runs.



HOW SKILLS DIFFER FROM SAVED PROMPTS

A saved prompt is a text snippet you paste into a chat. A Skill is an instruction set that ChatGPT executes — including step-by-step behavior, quality standards, an embedded template, and supporting resources.

A saved prompt tells you what to type. A Skill tells ChatGPT what to do.

Same Task. Different Results. A Skill Makes it Consistent.

Without a skill, you have to re-explain your requirements each time, which can lead to different results.

With a skill, you define the instructions once, and get consistent, higher-quality results.

Without a Skill

Same task, different results. You have to repeat instructions each time.

Attempt 1

Ask the model to create a project plan.



Create a project plan to launch a new mobile app. Include key steps and timeline.



1. Research the market
2. Design the app
3. Develop the app
4. Test the app
5. Launch the app

This should take about 6 months.

Attempt 2

Ask the model to create a project plan.



Create a project plan to launch a new mobile app. Include key steps and timeline.



1. Define requirements (2 weeks)
2. Build MVP (8 weeks)
3. User testing (4 weeks)
4. Marketing preparation (4 weeks)
5. Launch (2 weeks)

Total timeline: ~5 months.



Different structure, level of detail, and timeline. You have to rephrase and retry to get the result you want.

With a Skill

You define the instructions once. Get consistent, high-quality results.

1. Create the Skill

Define the task, style, and requirements once.



Project Plan Skill

- ✓ **Objective:** Create structured project plans
- ✓ **Audience:** Technical and non-technical stakeholders
- ✓ **Structure:** Phases, tasks, owners, timeline, and risks
- ✓ **Tone:** Clear, concise, professional
- ✓ **Format:** Use tables and bullet points



Skill saved.

2. Use the Skill (Anytime)

Run the same task and get consistent results.



Use my Project Plan skill to create a plan for launching a new mobile app.



Mobile App Launch Plan

Phase	Key Tasks	Timeline	Owner
1. Discovery	Market research, user interviews	2 weeks	Product
2. Planning	Requirements, technical design	3 weeks	Product / Eng
3. Development	Build MVP	8 weeks	Engineering
4. Testing	QA, user testing	4 weeks	QA
5. Launch	Marketing, release, monitoring	2 weeks	Product / Mktg



Consistent structure, format, and level of detail. You can use this skill anytime and get reliable results.




Skills turn one-time instructions into reusable, consistent workflows.

Less rework. Higher quality outputs. More time for what matters.

SKILL INSTRUCTIONS VS. SYSTEM PROMPTS

A system prompt like a project instruction establishes a persistent persona or context for a workspace. A Skill instruction describes a specific task process: inputs, steps, quality rules, and output format.

System prompts answer "who are you and what is your context." Skill instructions answer "how do you run this task from start to finish."



THE FIVE ELEMENTS OF A SKILL

Every well-built Skill contains five elements:

- **Opening** — what this Skill does and when to invoke it
- **Input specification** — exactly what the user must provide
- **Step-by-step process** — what ChatGPT does, in what order, producing what output
- **Quality standards** — rules that apply to every output regardless of input
- **Output template** — the exact structure of the final deliverable


Each element has a job. Missing one creates ambiguity that degrades output consistency across invocations.

WRITING THE OPENING AND INPUT SPECIFICATION

The opening tells users when and why to invoke the Skill. The input specification tells them exactly what to provide before it runs.

Vague inputs produce inconsistent outputs.

"Provide relevant details" is not an input specification. "Provide the competitor name, the service line to compare, and the date of analysis" is.




Anatomy of a Skill Markdown File

A skill uses a markdown file to provide the AI with instructions, tools, and context so it can complete specific tasks more reliably.

What is a Skill?

A skill is a reusable set of instructions, files, and optional tools that can be invoked across different chats to help the AI perform specialized tasks.

Typical Structure



Benefits

- ✓ Consistent, high-quality results
- ✓ Reusable across chats and projects
- ✓ Encapsulates best practices
- ✓ Reduces repetitive instructions
- ✓ Easier to share with others

```
1  ---
2  name: Data Analysis Assistant
3  description: Helps users analyze data, create visualizations,
4  and generate insights.
5  version: 1.0.0
6  author: N/A
7  tags: [data, analysis, visualization, python]
8  ---
9
10 # Purpose
11 Help users explore, analyze, and visualize data using Python.
12 Focus on clarity, accuracy, and actionable insights.
13
14 # Instructions
15 1. Understand the user's goal and the data they are working with.
16 2. Ask clarifying questions if the request is unclear.
17 3. Provide well-structured, efficient, and commented code.
18 4. Explain the results in simple terms.
19 5. Suggest next steps or additional analyses when relevant.
20
21 # Tools (optional)
22 - python: Use for data analysis, calculations, and visualizations.
23 - file_search: Use to find and read relevant files.
24 - web_search: Use for up-to-date information when needed.
25
26 # Knowledge / Context
27 - Prefer pandas, matplotlib, and seaborn for data analysis.
28 - Follow best practices for clean, reproducible code.
29 - Be mindful of data privacy and avoid exposing sensitive data.
30
31 # Examples
32 ## User Request
33 "Load a CSV file and show me a bar chart of sales by category."
34
35 ## Skill Output (shortened)
36 ```python
37 # Load data
38 import pandas as pd
39 ```
40
41 # Notes
42 - Keep explanations concise and practical.
43 - If data is large, suggest sampling or aggregation.
44 - Always verify results and call out limitations.
```

- 1 Metadata (Frontmatter)**
Basic information about the skill. Helps with organization, search, and discovery.
`name: Skill name`
`description: What it does`
`version: Version number`
`author: (optional)`
`tags: Relevant keywords`
- 2 Purpose**
High-level goal of the skill. Defines when and why to use it.
- 3 Instructions**
Step-by-step guidance for how the model should behave. Be clear and specific.
- 4 Tools (optional)**
Tools the skill can use, such as code execution, file search, or web search.
- 5 Knowledge / Context**
Important background information, domain rules, or reference material.
- 6 Examples**
Show example user requests and expected outputs. Helps the model follow your preferred format and style.
- 7 Notes**
Additional guidance, constraints, or best practices.



A well-designed skill leads to more accurate, consistent, and useful results.

Keep it clear, focused, and up to date. Customize the structure to fit your needs.

STEP-BY-STEP INSTRUCTIONS: PRECISION MATTERS


Skill steps must describe observable behavior. "Search for the competitor's publicly stated fee structure and note the retrieval date" is a step. "Gather information about the competitor" is not.

Vague verbs — gather, analyze, consider — leave interpretation to ChatGPT, which means different invocations produce different behavior. Replace them with observable actions.

AGENTS.md — Instructions for the Agent

AGENTS.md gives Codex persistent, repository- or directory-scoped instructions and context. It tells the agent how to work in this codebase, when to use skills, and what to follow.

Example Repository Structure



Directory Scoping and Precedence



```
1 # AGENTS.md
2 Instructions for Codex when working in this repository.
3
4 ## Repository expectations
5 - Follow the existing code style and conventions.
6 - Prefer simple, maintainable solutions.
7 - Do not introduce unnecessary dependencies.
8 - Consider security, performance, and accessibility.
9
10 ## How to work
11 - Understand the task and ask clarifying questions if needed.
12 - Read relevant docs before making changes (see Project map).
13 - Make small, focused changes with clear commit messages.
14 - Run tests and linting after code changes.
15 - Update documentation when behavior or APIs change.
16
17 ## Project map
18 - Architecture: docs/architecture.md
19 - API reference: docs/api.md
20 - Frontend: src/app/
21 - Tests: src/tests/
22 - Deployment: docs/deployment.md
23
24 ## Skills
25 - Use the data analysis skill for data exploration and reporting
26   (.agents/skills/data-analysis/SKILL.md).
27 - Use other skills in .agents/skills/ when relevant.
28 - Do not attempt complex workflows manually if a skill exists.
29
30 ## Validation
31 - Run 'npm test' and 'npm run lint' before finishing.
32 - Ensure all tests pass.
33 - Verify changes follow project conventions.
34 - Summarize what you changed and next steps.
```

AGENTS.md = how to operate here.

SKILL.md = how to perform a specialized repeatable workflow.

- Persistent Instructions**
AGENTS.md provides persistent, repository- or directory-scoped instructions that Codex reads before working. It tells the agent how to operate in this codebase.
- Project Context**
Define repository conventions, common commands, constraints, and links to architecture or other important information.
- Workspace Map**
Point to deeper documentation (e.g., docs/architecture.md) instead of putting everything in AGENTS.md.
- Skill Guidance**
Tell the agent when and how to use relevant skills. The SKILL.md file contains the reusable, specialized workflow with its own instructions, tools, and templates.
- Validation Rules**
Specify how to verify the work — tests, linting, code style, or other checks — before the agent considers the task complete.

How It Works

Codex follows your instructions, but you stay in control.



1. User Request

You give a task to Codex in this repository.



2. Reads AGENTS.md

Codex reads applicable AGENTS.md files (global, repo, and current directory).



3. Selects Skill (if needed)

Codex may load a relevant skill when appropriate. Skills are not executed automatically.



4. Performs Work

Uses repository files, tools, and skills (if chosen) to complete the task.



5. Validates Result

Runs checks (e.g., tests, lint), reviews the changes, and provides a summary.


Important Notes:

- AGENTS.md does not grant new tool permissions.
- Codex does not automatically run skills — it follows your instructions and chooses when to use them.
- Keep AGENTS.md focused, up to date, and reference deeper docs for details.

QUALITY STANDARDS: ENCODING TEAM KNOWLEDGE

Quality standards are the rules that apply to every output the Skill produces, regardless of input: always cite sources, flag information older than 12 months, use "client" not "customer," avoid speculation presented as fact.

Without explicit standards, different invocations of the same Skill will reflect different assumptions about what "good" means.



THE OUTPUT TEMPLATE


An output template defines the exact sections, order, and structure of the final deliverable. It is the contract between the Skill and anyone who reads its output.

Write placeholder text in each section to show what belongs there. A template makes the expected output unmistakable and makes the Skill testable against a clear standard.

TESTING FOR EDGE CASES IN SKILL INVOCATION

After testing the Skill with typical inputs, test the edges: a competitor with very little public information, a focus area that only partially applies, inputs that are slightly ambiguous.

Edge cases reveal where the Skill instructions leave too much to inference. Fix those gaps before sharing the Skill with your team.



WHEN TO USE A SKILL VS. A PROJECT VS. A CUSTOM GPT

Use a Project when work is ongoing and accumulates over time. Use a Skill when a task is discrete, recurring, and needs to be portable across projects. Use a custom GPT when you need a persistent AI persona with its own knowledge base that others can access.

Each tool solves a different problem. The mistake is treating whichever tool you learned first as the solution to every problem.

Three Ways to Work Smarter with AI

Different tools. Different strengths. A place for each.



Project

YOUR ONGOING WORKSPACE



Persistent workspace

Stays with you, accumulates over time.



Scoped to one initiative

Keep everything for a specific goal in one place.



Instructions + files + conversations

Bring together your guidance, reference materials, and ongoing discussions.



Skill

YOUR PORTABLE EXPERTISE



Portable

Take it with you, invocable anywhere.



Best for discrete recurring tasks

Save time on repeatable work.



Instructions + template + standards

Define how it works, what to produce, and the quality bar.



Custom GPT

YOUR AI PERSONA



Persistent persona

Maintains its identity, style, and behavior over time.



Shareable externally

Give others access, inside or outside your organization.



Own knowledge base

Use your content to provide specialized, reliable answers.



Suited to external audiences

Ideal for customers, partners, communities, or the public.

SAME AI. MORE POSSIBILITIES.

IDEAS → SKILLS → PROJECTS → IMPACT

SKILL MAINTENANCE OVER TIME

Skills go stale. Quality standards change, team terminology evolves, output templates get revised. A Skill that was accurate six months ago may produce outdated output today.

Assign someone to review each Skill quarterly. Update standards and templates before running the Skill on a high-stakes deliverable.

THE SKILL LIBRARY: MANAGING SKILLS AT THE TEAM LEVEL

When a team has multiple Skills, treat them as a library: document what each Skill does, who owns it, and when it was last reviewed. Skills without owners drift — no one updates them when standards change.

Ownership is not about permissions. It is about accountability for keeping the Skill accurate.



RESEARCH → SYNTHESIS → DRAFT → CRITIQUE: THE PATTERN


Multi-stage workflows consistently outperform single-prompt approaches for complex deliverables. Each stage produces a specific intermediate artifact that is reviewed before the next stage begins.

The pattern is not a preference — it is the difference between output you can defend and output you hope is right.

WHY PHASES BEAT ALL-IN-ONE PROMPTS

Asking ChatGPT to research, synthesize, draft, and critique in a single prompt produces a response that does all four things superficially. Running each phase separately, with review between, produces output where each phase is done well.

Intermediate outputs are not just stopping points — they are what you review to catch problems before they compound.



THE RESEARCH PHASE: RAW MATERIAL FIRST


In the research phase, ask for facts and sources only — not prose, not a draft, not analysis. The output should be a bullet-point findings list with citations.

Raw material is easier to evaluate than synthesized output. You can assess each claim individually and decide what to include before any drafting begins.

SYNTHESIS: STRUCTURE BEFORE YOU WRITE

After reviewing research findings, ask ChatGPT to group and structure the material into themes. Produce an outline — not prose, not the document.

The approved outline becomes the contract for the draft phase. If the outline is wrong, the draft will be wrong. Fix it before writing begins.



THE GENERATION PHASE: PRECISE CONSTRAINTS


When issuing the generation prompt, be explicit: word count, section structure, tone, audience, and any content the draft must include. Do not leave format to inference.

Constraints are specifications, not limitations. The more precise the specification, the closer the first draft will be to the final deliverable.

AI SELF-CRITIQUE: MAKING IT SPECIFIC

Generic critique prompts produce generic critique. Ask ChatGPT to identify specific problems: claims beyond what the research supports, "so what" statements that are generic rather than tailored to your context, sections that could be cut, and the weakest paragraph with an explanation of why.

Specific critique prompts surface real issues. "What could be improved?" does not.



HUMAN REVIEW: WHAT YOU DEFEND IS WHAT YOU OWN


After the critique phase, every sentence that remains in the document is one you could defend if your manager asked where it came from. If you cannot explain a sentence's origin, it either needs a source or needs to be cut.

AI-generated content becomes your work product when you approve it. Review accordingly.

THE QUARTERLY WORKFLOW: DESIGNING FOR REPETITION

A quarterly deliverable produced by an ad hoc process gives inconsistent results each quarter. A quarterly deliverable run through a documented workflow gives consistent results that improve over time as the workflow is refined.

The first time you run the workflow is the hardest. After that, you are refining rather than rebuilding.



HANDLING CONTRADICTIONARY RESEARCH FINDINGS

When sources disagree, do not let ChatGPT average them. Note the contradiction explicitly and investigate the source: different time periods, different markets, or different methodologies.

Contradictions are information. A finding that two authoritative sources disagree is often more valuable than a fabricated consensus between them.

From Research to Refined Result

A collaborative workflow between AI and humans

● AI step

● Human checkpoint



BETTER INFORMATION. BETTER DECISIONS.

MAKING WORKFLOWS REPEATABLE AND TRANSFERABLE

A workflow that only you can run is a personal process, not an organizational asset. Repeatable workflows include prompt templates with [PLACEHOLDER] syntax, defined checkpoints, output specifications, and examples of acceptable output.

The document is what you preserve. The deliverable it produces is disposable after the quarter ends.

LAB 6



POP QUIZ: In a multi-stage workflow, what is the correct output of the research phase?

- A. A full draft of the deliverable written based on initial findings.
- B. A structured outline with section headings and sub-bullets.
- C. A bullet-point findings list with sources only.
- D. A critique of the research approach used.



POP QUIZ: In a multi-stage workflow, what is the correct output of the research phase?

- A. A full draft of the deliverable written based on initial findings.
- B. A structured outline with section headings and sub-bullets.
- C. A bullet-point findings list with sources only.**
- D. A critique of the research approach used.



POP QUIZ: When asking ChatGPT to critique a draft, which approach produces the most useful feedback?

- A. "What could be better?" — an open question that allows full creative latitude.
- B. "Identify claims beyond research support, generic statements, and the weakest paragraph."
- C. "Rewrite the entire document to be stronger."
- D. "Give me a quality score from 1 to 10."



POP QUIZ: When asking ChatGPT to critique a draft, which approach produces the most useful feedback?


- A. "What could be better?" — an open question that allows full creative latitude.
- B. "Identify claims beyond research support, generic statements, and the weakest paragraph."**
- C. "Rewrite the entire document to be stronger."
- D. "Give me a quality score from 1 to 10."



WORKFLOW SPECIFICATION: THE SOURCE OF TRUTH

A workflow specification answers five questions before any prompt is written: what are the inputs, what are the discrete steps, what does each step produce, where are the human checkpoints, and what does the final deliverable look like?

The specification is what you share with a colleague. The prompts implement the specification.



DEFINING INPUTS BEFORE YOU PROMPT

Inputs are the variable parts of a workflow — the competitor name, the date range, the business unit, the topic focus. Defining them explicitly means users know exactly what to provide before starting.

Implicit inputs produce inconsistent outputs. Different analysts supply different things and get different results from the same workflow.

WRITING PROMPT TEMPLATES WITH PLACEHOLDERS

Use [PLACEHOLDER] syntax to mark variable parts of each prompt. A template like "Research [COMPETITOR NAME]'s publicly stated fee structure as of [DATE]" tells the user exactly what to substitute.

Templates are the difference between a workflow a new analyst can follow and one they have to reverse-engineer by asking the analyst who wrote it.



CALIBRATING PROMPT SPECIFICITY


A prompt that is too vague produces inconsistent output. A prompt that is too specific may not accommodate natural variation across workflow runs.

Calibration comes from testing: run the workflow twice with different inputs and note where outputs diverge unexpectedly. That divergence points to where the prompt needs tightening.

TESTING AND REFINING WORKFLOW STEPS

After building a workflow, run it end-to-end before sharing it. Note every step that required a follow-up prompt to produce the right output. That follow-up belongs in the step's template.

A step that needed three clarifying turns is a step whose template needs to be rewritten.



DOCUMENTING FOR TEAM HANDOFF


A workflow document should contain enough information for a team member who was not in this training to run it successfully: prompt templates, checkpoint instructions, output specifications, and examples of acceptable output.

The test is simple: give it to someone who has never seen it and watch where they get stuck. Fix every sticking point before you share it.

THE WORKFLOW DOCUMENT VS. THE DELIVERABLE

The workflow document is the asset that compounds in value over time. The deliverable it produces is current for one quarter and then disposable.

Invest in the document. The deliverable is a byproduct — the document is what you will still be using two years from now.



LAB 7



POP QUIZ: A workflow step says "gather information about the competitor." What is wrong with this instruction?

- A. It is too long and will confuse ChatGPT with unnecessary detail.
- B. It does not specify what information, in what format, or from what sources.
- C. It uses "competitor" without naming a specific competitor in the instruction.
- D. It should be split into three separate workflow documents.



POP QUIZ: A workflow step says "gather information about the competitor." What is wrong with this instruction?

- A. It is too long and will confuse ChatGPT with unnecessary detail.
- B. It does not specify what information, in what format, or from what sources.**
- C. It uses "competitor" without naming a specific competitor in the instruction.
- D. It should be split into three separate workflow documents.



POP QUIZ: Where must a human checkpoint exist in a reusable workflow?

- A. Only at the end, after the final deliverable is produced.
- B. At every step, to review each intermediate output before continuing.
- C. At points where errors compound — before raw material becomes synthesis, and before the draft is approved.
- D. Human checkpoints are optional if the workflow has been run successfully before.



POP QUIZ: Where must a human checkpoint exist in a reusable workflow?


- A. Only at the end, after the final deliverable is produced.
- B. At every step, to review each intermediate output before continuing.
- C. At points where errors compound — before raw material becomes synthesis, and before the draft is approved.**
- D. Human checkpoints are optional if the workflow has been run successfully before.



FROM WORKFLOW TO SKILL: THE NEXT STEP

A reusable workflow document is a set of instructions a human follows. A Skill is a set of instructions ChatGPT follows when invoked. The Skill encodes the workflow so the model executes the steps rather than the analyst coordinating each prompt manually.

The Skill eliminates manual prompt coordination — the behavior is built in.



PLANNING A SKILL BEFORE WRITING INSTRUCTIONS


The planning table comes before the Skill interface. Define the Skill name, trigger description, required inputs, step sequence, output format, quality standards, and embedded template before writing a word of instructions.

Writing without planning produces instructions that are too vague to enforce or too rigid to use in practice.

AVOIDING VAGUE INSTRUCTION VERBS

The most common Skill writing mistake is using verbs that do not specify behavior: gather, analyze, consider, review. These leave interpretation to ChatGPT, which means different invocations produce different behavior.

Replace every vague verb with an observable action: "search for," "produce a table comparing," "write a one-paragraph summary of the top three differences."



WRITING QUALITY STANDARDS THAT SURVIVE TURNOVER


Quality standards should be written so they are understandable to someone who has never worked at your organization. Do not rely on context that only long-tenured employees have.

A standard that says "follow the usual format" is not a standard. One that says "use the four-section memo format: Summary, Findings, Recommendation, Next Steps" can be followed by anyone from day one.

SUPPORTING RESOURCES INSIDE A SKILL

Skills can include reference materials that ChatGPT uses during execution: a sample completed report showing expected quality and depth, a terminology glossary with org-specific terms and abbreviations, and a list of approved sources for the task type.

Supporting resources remove ambiguity about what "good" looks like. Anyone invoking the Skill starts with the same context.



SKILL INVOCATION: GIVING THE RIGHT INPUTS THE FIRST TIME

When invoking a Skill, provide all required inputs in the first message. Missing inputs lead to clarifying back-and-forth that breaks the Skill's designed flow.

If you find yourself frequently supplementing inputs after invocation, the input specification in the Skill instructions needs to be more explicit about what is required.

TESTING A SKILL FROM A FRESH CONVERSATION

The correct test for a Skill is to invoke it in a conversation with no prior context — no project instructions, no earlier turns establishing background. If the Skill requires outside knowledge to produce correct output, it is incomplete.

Test it with only the inputs it asks for and evaluate the output against the quality checklist you defined.

QUALITY CONTROL CHECKLIST			
ITEM	TOLERANCES	ACCEPTANCE CRITERIA	SIGNATURES

INSPECTED BY _____
DATE _____
APPROVED BY _____

APPROVED BY _____

SKILLS VS. WORKFLOW DOCUMENTS: WHAT TO USE WHEN

A workflow document is best when the analyst must exercise judgment at each step and the task varies significantly by input. A Skill is best when the task is highly structured, the output template is fixed, and consistency matters more than flexibility.

Both encode organizational knowledge. One runs on a person; the other runs on the model.

BUILDING A CULTURE OF REUSABLE AI WORK

Workflows and Skills only create organizational value if they are shared, maintained, and used beyond the person who built them. This requires two things: a place to store them where teammates can find them, and a norm that building reusable tools is part of the job.

The first person who hands a working workflow to a colleague creates the norm. That is usually the person who attended the training.

LAB 8



POP QUIZ: What is the primary advantage of a Skill over a saved prompt stored in a text file?

- A. Skills are free; saved prompts require a subscription upgrade.
- B. Skills encode step-by-step behavior, quality standards, and output templates that ChatGPT executes automatically.
- C. Skills are longer, which makes them inherently more detailed than a saved prompt.
- D. Skills are cloud-stored, so they are safer from accidental deletion.



POP QUIZ: What is the primary advantage of a Skill over a saved prompt stored in a text file?

- A. Skills are free; saved prompts require a subscription upgrade.
- B. Skills encode step-by-step behavior, quality standards, and output templates that ChatGPT executes automatically.**
- C. Skills are longer, which makes them inherently more detailed than a saved prompt.
- D. Skills are cloud-stored, so they are safer from accidental deletion.



POP QUIZ: An analyst writes a Skill step that says "analyze the competitor's digital experience." What should they write instead?

- A. "Consider the competitor's digital experience holistically and note your impressions."
- B. "Gather digital experience information from any available public sources."
- C. "Search for user reviews and ratings of the competitor's mobile app and website on credible review platforms and produce a three-bullet summary."
- D. "Review the competitor's digital experience and note any important findings you observe."



POP QUIZ: An analyst writes a Skill step that says "analyze the competitor's digital experience." What should they write instead?


- A. "Consider the competitor's digital experience holistically and note your impressions."
- B. "Gather digital experience information from any available public sources."
- C. "Search for user reviews and ratings of the competitor's mobile app and website on credible review platforms and produce a three-bullet summary."**
- D. "Review the competitor's digital experience and note any important findings you observe."



KEY PRINCIPLES FOR AI-POWERED WORKFLOWS

Three principles hold across every workflow and Skill covered today:

- **Design before you prompt** — the specification comes first; the prompt implements it
- **Human checkpoints are not optional** — they are the accountability structure that makes AI output defensible
- **The document is the asset** — workflows persist; deliverables are disposable




DAY 2 KEY TAKEAWAYS

- One-off chats are disposable. Workflows and Skills are organizational assets that compound in value.
- Connected applications change the data risk profile — plan permissions deliberately before enabling access.
- Projects persist context across conversations; Skills port that context to any conversation.
- Multi-stage workflows produce more defensible output than single-prompt approaches.
- The workflow document is what you preserve. The deliverable is what you present.

COMING TOMORROW: DAY 3

Day 3 introduces Codex and AI-assisted software development. You will explore an unfamiliar codebase, find and fix a bug, add a feature, generate tests, and conduct a code review — all through natural language.

No programming experience is required. The goal is to supervise AI-generated code, not to write it yourself.



CONGRATS!



Today covered turning ChatGPT use into repeatable team workflows.

Connected apps: data risk, permissions, and deliberate access decisions.

Projects and Skills: persistent context, portable instructions, and testing.

Multi-stage workflows: specifications, staged outputs, and human checkpoints.

The workflow document is the asset; the deliverable is disposable.